Multiple linear regression assumptions and multicollinearity

In prior videos, you have learned about linear regression assumptions. In this reading, you will build off that knowledge base to extend your understanding of multiple linear regression assumptions. This reading will help you review assumptions that apply to both simple linear regression and multiple linear regression, and will then focus more heavily on the concept of multicollinearity.

Recall that simple linear regression has four main assumptions that provide validity to the results derived from the analysis. To this list of four assumptions, we add the no multicollinearity assumption when working with multiple linear regression.

1. **Linearity:** Each predictor variable (*Xi*) is linearly related to the outcome variable (Y).

2. **(Multivariate) normality:** The errors are normally distributed.*

3. **Independent observations:** Each observation in the dataset is independent.

4. **Homoscedasticity:** The variation of the errors is constant or similar across the model.*

5. **No multicollinearity:** No two independent variables (*Xi* and *Xj*) can be highly correlated with each other.

*Note on errors and residuals

As noted earlier, “residuals” and “errors” are sometimes used interchangeably, but there is a difference. We use residuals to estimate errors when we are checking the normality and homoscedasticity assumptions of linear regression.

- **Residuals** are the difference between the predicted and observed values. You can calculate residuals after you build a regression model by subtracting the predicted values from the observed values.

- **Errors** are the natural noise assumed to be in the model.

Extending prior assumptions

Much of what you learned about the first four assumptions with regard to simple linear regression can be directly applied to multiple linear regression. The code might be slightly different or longer, but the rationale is the same.

Linearity

- With multiple linear regression, you need to consider whether each *x* variable has a linear relationship with the *y* variable.

- You can make multiple scatterplots instead of just one, using seaborn's pairplot() function, or the scatterplot() function multiple times. Other libraries with plotting capabilities will have similar functions.

Independent observations

- The independent observations assumption is still primarily focused on data collection.

- You can check the validity of the assumption in the same way you would with simple linear regression.

(Multivariate) Normality

- Just as with simple linear regression, you can construct the model, and then create a Q-Q plot of the residuals.

- If you observe a straight diagonal line on the Q-Q plot, then you can proceed in your analysis. You can also plot a histogram of the residuals and check if you observe a normal distribution that way.

- **Note:** It’s a common misunderstanding that the independent and/or dependent variables must be normally distributed when performing linear regression. This is not the case. Only the model’s residuals are assumed to be normal.

Homoscedasticity

- As with simple linear regression, for multiple linear regression, just create a plot of the residuals vs. fitted values.

- If the data points seem to be scattered randomly across the line where residuals equal 0, then you can proceed.

How to check the no multicollinearity assumption

The no multicollinearity assumption is unique to multiple linear regression as it focuses on potential relationships between different independent (X) variables. When assessing the no multicollinearity assumption, you’re interested in identifying any linear relationships between the independent (X) variables. X variables that are linearly related could muddle the interpretation of the model’s results. If there are X variables that are linearly related, it is usually best to remove some independent variables from the model.

Note, however, that the assumption of no multicollinearity is most important when you are using your regression model to make inferences about your data, because the inclusion of collinear data increases the standard errors of the model's beta parameter estimates. But there may be times when the primary purpose of your model is to make predictions and the need for accurate predictions outweighs the need to make inferences about your data. In this case, including the collinear independent variables may be justified because it's possible that their inclusion would result in better predictions.

There are a few ways to check the no multicollinearity assumption. This reading will cover two of them. One is purely visual, and the other is numerical in nature. Both can be done prior to building the linear regression model.

Scatterplots or Scatterplot Matrix

A visual way to identify multicollinearity between independent (X) variables is using scatterplots or scatterplot matrices. The process is the same as when you checked the linearity assumption, except now you're just focusing on the X variables, not the relationship between the X variables and the Y variable. If you're using the seaborn library, you can use the *pairplot* function, or the *scatterplot* function multiple times.

Variance Inflation Factors (VIF)

Calculating the variance inflation factor, or VIF, for each independent (X) variable is a way to quantify how much the variance of each variable is "inflated" due to correlation with other X variables. You can read more about VIFs on the Pennsylvania State University's Eberly College of Science website or on the website for Vilnius University's e-book on *Practical Econometrics and Data Science*. The details of calculating VIF are beyond the scope of this course, but it's helpful to know that *VIFi* represents the amount that the standard error of coefficient *βi* increases relative to a situation in which all of the predictor variables are uncorrelated.

To calculate the VIF for each predictor variable, you can use the variance_inflation_factor() function from the *statsmodels* package. Here is an example of how you might obtain VIFs for your predictor variables.

The smallest value a VIF can take on is 1, which would indicate 0 correlation between the X variable in question and the other predictor variables in the model. A high VIF, such as 5 and above, according to the statsmodels documentation, can indicate the presence of multicollinearity.

What to do if there is multicollinearity in

your model

Variable Selection

The easiest way to handle multicollinearity is simply to only use a subset of independent variables in your model. For example, if your multiple linear regression model is something like this:

$y=\beta 0+\beta 1X1+\beta 2X2+\beta 3X3$

But if $X1$ and $X3$ are highly correlated, then you can choose to include only $X1$ or $X3$ in your final model, but not both.

There are a few specific statistical techniques you can use to select variables strategically. You'll learn about these more in future videos:

- Forward selection
- Backward elimination

Advanced Techniques

In addition to the techniques listed above that will be covered in-depth in this course, there are more advanced techniques that you may come across in your career as a data professional, such as:

- Ridge regression
- Lasso regression
- Principal component analysis (PCA)

These techniques can result in more accurate and predictive models, but can complicate the interpretation of regression results.

Key Takeaways

- Many of the assumptions of simple linear regression extend readily to multiple linear regression.
- You can use scatterplots and variance inflation factors to check for multicollinearity in a regression model.
- There are different techniques for variable selection to remove multicollinearity in a model.